

String Manipulation

Dav King

4/12/22

Main Ideas

- Working with string data is essential for a number of data science tasks, including data cleaning, data preparation, and text analysis.
- The **stringr** package in R (part of the **tidyverse**) contains useful tools for working with character strings.

Coming Up

- HW due on Thursday at 11:59 PM
- Lab 2 in groups on Thursday

Lecture Notes and Exercises

In addition to the **tidyverse**, we will use the **stringr** package.

```
library(tidyverse)
library(stringr)
```

stringr provides tools to work with character strings. Functions in **stringr** have consistent, memorable names.

- All begin with **str_** (**str_count()**, **str_detect()**, **str_trim()**, etc).
- All take a vector of strings as their first arguments.
- We only have time to explore the basics. I encourage you to explore on your own using the **additional resources** below.

Preliminaries

Character strings in R are defined by double quotation marks. These can include numbers, letters, punctuation, whitespace, etc.

```
string1 <- "STA 199 is my favorite class"
string1
```

```
## [1] "STA 199 is my favorite class"
```

You can combine character strings in a vector.

```
string2 <- c("STA 199", "Data Science", "Duke")
string2
```

```
## [1] "STA 199"      "Data Science" "Duke"
```

Question: What if we want to include a quotation in a string? Why doesn't the code below work?

You interrupt the parentheses of the string.

```
string3 <- "I said "Hello" to my class"
```

To include a double quote in a string **escape it** using a backslash. Try it now in the code chunk below and name your string `string4`.

```
string4 <- "I said \"Hello\" to my class"
string4
```

```
## [1] "I said \"Hello\" to my class"
```

If you want to include an actual backslash, **escape it** as shown below. This may seem tedious but it will be important later.

```
string5 <- "\\\"
string5
```

```
## [1] "\\\"
```

The function `writeLines()` shows the content of the strings not including escapes. Try it for `string1`, `string2`, `string3`, `string4`, and `string5` in the code chunk below.

```
writeLines(c(string1, string2, string4, string5))
```

```
## STA 199 is my favorite class
## STA 199
## Data Science
## Duke
## I said "Hello" to my class
## \
```

U.S. States

To demonstrate the basic functions from `stringr` we will use a vector of all 50 U.S. states.

```
states <- c("alabama", "alaska", "arizona", "arkansas", "california",
            "colorado", "connecticut", "delaware", "florida", "georgia",
            "hawaii", "idaho", "illinois", "indiana", "iowa", "kansas",
            "kentucky", "louisiana", "maine", "maryland", "massachusetts",
            "michigan", "minnesota", "mississippi", "missouri", "montana",
            "nebraska", "nevada", "new hampshire", "new jersey",
```

```
"new mexico", "new york", "north carolina", "north dakota", "ohio",
"oklahoma", "oregon", "pennsylvania", "rhode island",
"south carolina", "south dakota", "tennessee", "texas", "utah",
"vermont", "virginia", "washington", "west virginia", "wisconsin",
"wyoming")
```

str_length() Given a string, return the number of characters.

```
string1
```

```
## [1] "STA 199 is my favorite class"
```

```
str_length(string1)
```

```
## [1] 28
```

Given a vector of strings, return the number of characters in each string.

```
str_length(states)
```

```
## [1] 7 6 7 8 10 8 11 8 7 7 6 5 8 7 4 6 8 9 5 8 13 8 9 11 8
## [26] 7 8 6 13 10 10 8 14 12 4 8 6 12 12 14 12 9 5 4 7 8 10 13 9 7
```

str_c() Combine two (or more) strings.

```
str_c("STA 199", "is", "my", "favorite", "class")
```

```
## [1] "STA 199ismyfavoriteclass"
```

Use **sep** to specify how the strings are separated.

```
str_c("STA 199", "is", "my", "favorite", "class", sep = " ")
```

```
## [1] "STA 199 is my favorite class"
```

str_to_lower() and **str_to_upper()**

Convert the case of a string from lower to upper or vice versa.

```
str_to_upper(states)
```

```
## [1] "ALABAMA"      "ALASKA"        "ARIZONA"       "ARKANSAS"
## [5] "CALIFORNIA"   "COLORADO"      "CONNECTICUT"   "DELAWARE"
## [9] "FLORIDA"      "GEORGIA"       "HAWAII"        "IDAHO"
## [13] "ILLINOIS"     "INDIANA"       "IOWA"          "KANSAS"
## [17] "KENTUCKY"     "LOUISIANA"     "MAINE"         "MARYLAND"
## [21] "MASSACHUSETTS" "MICHIGAN"      "MINNESOTA"     "MISSISSIPPI"
```

```
## [25] "MISSOURI"      "MONTANA"      "NEBRASKA"     "NEVADA"
## [29] "NEW HAMPSHIRE" "NEW JERSEY"   "NEW MEXICO"   "NEW YORK"
## [33] "NORTH CAROLINA" "NORTH DAKOTA" "OHIO"         "OKLAHOMA"
## [37] "OREGON"        "PENNSYLVANIA" "RHODE ISLAND" "SOUTH CAROLINA"
## [41] "SOUTH DAKOTA"   "TENNESSEE"    "TEXAS"        "UTAH"
## [45] "VERMONT"       "VIRGINIA"     "WASHINGTON"   "WEST VIRGINIA"
## [49] "WISCONSIN"     "WYOMING"
```

`str_sub()`

Extract parts of a string from `start` to `end`, inclusive.

```
str_sub(states, 1, 4)
```

```
## [1] "alab" "alas" "ariz" "arka" "cali" "colo" "conn" "dela" "flor" "geor"
## [11] "hawa" "idah" "illi" "indi" "iowa" "kans" "kent" "loui" "main" "mary"
## [21] "mass" "mich" "minn" "miss" "miss" "mont" "nebr" "neva" "new " "new "
## [31] "new " "new " "nort" "nort" "ohio" "okla" "oreg" "penn" "rhod" "sout"
## [41] "sout" "tenn" "texa" "utah" "verm" "virg" "wash" "west" "wisc" "wyom"
```

```
str_sub(states, -4, -1)
```

```
## [1] "bama" "aska" "zona" "nsas" "rnia" "rado" "icut" "ware" "rida" "rgia"
## [11] "waii" "daho" "nois" "iana" "iowa" "nsas" "ucky" "iana" "aine" "land"
## [21] "etts" "igan" "sota" "ippi" "ouri" "tana" "aska" "vada" "hire" "rsey"
## [31] "xico" "york" "lina" "kota" "ohio" "homa" "egon" "ania" "land" "lina"
## [41] "kota" "ssee" "exas" "utah" "mont" "inia" "gton" "inia" "nsin" "ming"
```

Practice: Combine `str_sub()` and `str_to_upper()` to capitalize each state.

```
str_sub(states, 1, 1) <- str_to_upper(str_sub(states, 1, 1))
states
```

```
## [1] "Alabama"      "Alaska"      "Arizona"     "Arkansas"
## [5] "California"   "Colorado"    "Connecticut" "Delaware"
## [9] "Florida"      "Georgia"     "Hawaii"      "Idaho"
## [13] "Illinois"     "Indiana"     "Iowa"        "Kansas"
## [17] "Kentucky"     "Louisiana"   "Maine"       "Maryland"
## [21] "Massachusetts" "Michigan"    "Minnesota"   "Mississippi"
## [25] "Missouri"     "Montana"     "Nebraska"    "Nevada"
## [29] "New hampshire" "New jersey"  "New mexico"  "New york"
## [33] "North carolina" "North dakota" "Ohio"        "Oklahoma"
## [37] "Oregon"       "Pennsylvania" "Rhode island" "South carolina"
## [41] "South dakota"  "Tennessee"   "Texas"       "Utah"
## [45] "Vermont"      "Virginia"    "Washington"  "West virginia"
## [49] "Wisconsin"    "Wyoming"
```

```
str_to_upper(states)
```

```
## [1] "ALABAMA"      "ALASKA"      "ARIZONA"     "ARKANSAS"
## [5] "CALIFORNIA"   "COLORADO"    "CONNECTICUT" "DELAWARE"
## [9] "FLORIDA"      "GEORGIA"     "HAWAII"      "IDAHO"
## [13] "ILLINOIS"     "INDIANA"     "IOWA"        "KANSAS"
## [17] "KENTUCKY"     "LOUISIANA"   "MAINE"       "MARYLAND"
## [21] "MASSACHUSETTS" "MICHIGAN"    "MINNESOTA"   "MISSISSIPPI"
## [25] "MISSOURI"     "MONTANA"     "NEBRASKA"    "NEVADA"
## [29] "NEW HAMPSHIRE" "NEW JERSEY"  "NEW MEXICO"  "NEW YORK"
## [33] "NORTH CAROLINA" "NORTH DAKOTA" "OHIO"        "OKLAHOMA"
## [37] "OREGON"       "PENNSYLVANIA" "RHODE ISLAND" "SOUTH CAROLINA"
## [41] "SOUTH DAKOTA"  "TENNESSEE"   "TEXAS"       "UTAH"
## [45] "VERMONT"      "VIRGINIA"    "WASHINGTON"  "WEST VIRGINIA"
## [49] "WISCONSIN"    "WYOMING"
```

`str_sort()` Sort a string. Below we sort in decreasing alphabetical order.

```
str_sort(states, decreasing = TRUE)
```

```
## [1] "Wyoming"      "Wisconsin"   "West virginia" "Washington"
## [5] "Virginia"     "Vermont"     "Utah"         "Texas"
## [9] "Tennessee"    "South dakota" "South carolina" "Rhode island"
## [13] "Pennsylvania" "Oregon"      "Oklahoma"     "Ohio"
## [17] "North dakota" "North carolina" "New york"     "New mexico"
## [21] "New jersey"   "New hampshire" "Nevada"       "Nebraska"
## [25] "Montana"      "Missouri"    "Mississippi"  "Minnesota"
## [29] "Michigan"     "Massachusetts" "Maryland"     "Maine"
## [33] "Louisiana"    "Kentucky"    "Kansas"       "Iowa"
## [37] "Indiana"      "Illinois"    "Idaho"        "Hawaii"
## [41] "Georgia"      "Florida"     "Delaware"     "Connecticut"
## [45] "Colorado"     "California"  "Arkansas"     "Arizona"
## [49] "Alaska"      "Alabama"
```

Regular Expressions

A **regular expression** is a sequence of characters that allows you to describe string patterns. We use them to search for patterns.

Examples of usage include the following data science tasks:

- extract a phone number from text data
- determine if an email address is valid
- determine if a password has some specified number of letters, characters, numbers, etc
- count the number of times “statistics” occurs in a corpus of text

To demonstrate regular expressions, we will use a vector of the states bordering North Carolina.

```
nc_states <- c("North Carolina", "South Carolina", "Virginia", "Tennessee",
              "Georgia")
```

Basic Match We can match exactly using a **basic match**.

```
str_view_all(nc_states, "in")
```

We can match any character using `.`

```
str_view_all(nc_states, ".a")
```

Question: What if we want to match a period `.`?

Escape it using `\.`

Another example using escapes:

```
str_view(c("a.c", "abc", "def"), "a\\.c")
```

Anchors Match the start of a string using `^`.

```
str_view(nc_states, "^G")
```

Match the end of a string using `$`.

```
str_view(nc_states, "a$")
```

str_detect() Determine if a character vector matches a pattern.

```
nc_states
```

```
## [1] "North Carolina" "South Carolina" "Virginia"      "Tennessee"
## [5] "Georgia"
```

```
str_detect(nc_states, "a")
```

```
## [1] TRUE TRUE TRUE FALSE TRUE
```

```
nc_states
```

```
str_subset()
```

```
## [1] "North Carolina" "South Carolina" "Virginia"      "Tennessee"
## [5] "Georgia"
```

```
str_subset(nc_states, "e$")
```

```
## [1] "Tennessee"
```

str_count() Determine how many matches there are in a string.

```
nc_states
```

```
## [1] "North Carolina" "South Carolina" "Virginia"      "Tennessee"  
## [5] "Georgia"
```

```
str_count(nc_states, "a")
```

```
## [1] 2 2 1 0 1
```

str_replace() and str_replace_all() Replace matches with new strings.

```
str_replace(nc_states, "a", "-")
```

```
## [1] "North C-rolina" "South C-rolina" "Virgini-"      "Tennessee"  
## [5] "Georgi-"
```

Use **str_replace_all()** to replace all matches with new strings.

```
str_replace_all(nc_states, "a", "-")
```

```
## [1] "North C-rolin-" "South C-rolin-" "Virgini-"      "Tennessee"  
## [5] "Georgi-"
```

Many Matches The regular expressions below match more than one character.

- Match any digit using `\d` or `[[:digit:]]`
- Match any whitespace using `\s` or `[[:space:]]`
- Match f, g, or h using `[fgh]`
- Match anything but f, g, or h using `[^fgh]`
- Match lower-case letters using `[a-z]` or `[[:lower:]]`
- Match upper-case letters using `[A-Z]` or `[[:upper:]]`
- Match alphabetic characters using `[A-Za-z]` or `[[:alpha:]]`

Remember these are regular expressions! To match digits you'll need to escape the `.`, so use `"\d"`, not `"."`

Practice

To practice manipulating strings we will use question and answer data from two recent seasons (2008 - 2009) of the television game show *Jeopardy!*.

```
jeopardy <- read_csv("questions.csv")
```

- **category:** category of question
- **value:** value of question in dollars
- **question:** text of question
- **answer:** text of question answer
- **year:** year episode aired

```
glimpse(jeopardy)
```

```
## Rows: 40,865
## Columns: 5
## $ category <chr> "OLD FOLKS IN THEIR 30s", "MOVIES & TV", "A STATE OF COLLEGE--
## $ value      <dbl> 200, 200, 200, 200, 200, 200, 400, 400, 400, 400, 400, 400, 6~
## $ question  <chr> "goop.com is a lifestyles website from this Oscar-winning act~
## $ answer    <chr> "Gwyneth Paltrow", "Jay Leno", "Texas", "a pride", "a bunny h~
## $ year      <dbl> 2009, 2009, 2009, 2009, 2009, 2009, 2009, 2009, 2009, 2009, 2~
```

- (1) Use a single code pipeline and a function from `stringr` to return all rows where the answer **contains** the word “Durham”

```
jeopardy[str_detect(jeopardy$answer, "Durham"),]
```

```
## # A tibble: 3 x 5
##   category      value question      answer year
##   <chr>         <dbl> <chr>      <chr>  <dbl>
## 1 BULL          2000 "\"Bull City\"", this place's nickname, is ~ Durham 2009
## 2 BASEBRAWL     1000 "In 1995 10 players were ejected for a bra~ the D~ 2009
## 3 MOVIES BY QUOTE 800 "Crash: \"Man, that ball got out of here i~ Bull ~ 2009
```

- (2) Use a single code pipeline and `stringr` to find the length of all of the answers, sort by decreasing length, and return the five longest answers.

```
jeopardy %>%
  mutate(ans_length = str_length(answer)) %>%
  arrange(desc(ans_length)) %>%
  slice(1:5) %>%
  select(answer, ans_length)
```

```
## # A tibble: 5 x 2
##   answer      ans_length
##   <chr>         <int>
## 1 a microphone & the masks of comedy & tragedy (a TV set, a movie ca~      86
## 2 hiding your light under a bushel (keeping your light underneath a ~      82
## 3 International Talk Like a Pirate Day (National Talk Like a Pirate ~      79
## 4 (any of) the (St. Louis) Rams, the Oakland Raiders, or the San Die~      77
## 5 to take the number that's between 3 and 5 (averaging the 2 middle ~      74
```

- (3) What answer has the most digits?

```
jeopardy %>%
  mutate(digits = str_count(answer, "[0-9]")) %>%
  arrange(desc(digits)) %>%
  slice(1) %>%
  select(answer)
```

```
## # A tibble: 1 x 1
##   answer
##   <chr>
## 1 1939 (or 1942)
```


(4) Return all rows where the category has a period.

```
jeopardy[str_detect(jeopardy$category, "\\."),]
```

```
## # A tibble: 1,249 x 5
##   category          value question          answer year
##   <chr>          <dbl> <chr>          <chr> <dbl>
## 1 I LOVE L.A. KERS      400 "Kobe called it \"idiotic criticism\" ~ Shaqu~ 2009
## 2 I LOVE L.A. KERS      800 "A wizard at passing the ball, this La~ Magic~ 2009
## 3 I LOVE L.A. KERS     1200 "This Laker giant was nicknamed \"The ~ Wilt ~ 2009
## 4 I LOVE L.A. KERS     1600 "This Hall-of-Fame guard & former Lake~ Jerry~ 2009
## 5 I LOVE L.A. KERS     2000 "This flashy Lakers forward was nickna~ James~ 2009
## 6 IT'S AN L.A. THING     200 "Wanna live in this city, 90210? in Ju~ Bever~ 2009
## 7 IT'S AN L.A. THING     400 "Originally the letters in this landma~ the H~ 2009
## 8 IT'S AN L.A. THING     600 "Good times are Bruin in this district~ Westw~ 2009
## 9 IT'S AN L.A. THING     800 "You can hit the Comedy Store, House o~ Sunse~ 2009
## 10 IT'S AN L.A. THING   1000 "Originally called \"Nuestro Pueblo\" ~ the W~ 2009
## # ... with 1,239 more rows
```

(5) Using a single code pipeline, return all rows where the question contains a (numeric) year between 1800 and 1999

```
jeopardy %>%
  mutate(year = str_detect(jeopardy$question, paste(c("18\\d\\d", "19\\d\\d"),
                                                    collapse = "|"))) %>%
  filter(year == T)
```

```
## # A tibble: 6,749 x 5
##   category          value question          answer year
##   <chr>          <dbl> <chr>          <chr> <lgl>
## 1 AMERICAN AUTHORS      800 "During the War~ Washi~ TRUE
## 2 MATHEM-ATTACK!     1200 "(<a href=\"htt~ a mat~ TRUE
## 3 AMERICAN AUTHORS     2000 "He reviewed fi~ Phili~ TRUE
## 4 AMERICAN AUTHORS      200 "While he was i~ Hemin~ TRUE
## 5 AMERICAN AUTHORS      400 "In 1884 she mo~ Willa~ TRUE
## 6 BEST PICTURE OSCAR-WINNERS IN OTHER WORDS 400 "1980: \"Regula~ Ordin~ TRUE
## 7 DOWN MEXICO WAY      400 "In 1986 Mexico~ the W~ TRUE
## 8 BEST PICTURE OSCAR-WINNERS IN OTHER WORDS 800 "1932: \"Magnif~ Grand~ TRUE
## 9 BEST PICTURE OSCAR-WINNERS IN OTHER WORDS 1200 "1976: \"A Sing~ Rocky TRUE
## 10 BEST PICTURE OSCAR-WINNERS IN OTHER WORDS 1600 "1954: \"Docksi~ On th~ TRUE
## # ... with 6,739 more rows
```

(6) Using a single code pipeline, return all rows with answers that begin with three vowels.

```
jeopardy %>%
  mutate(vowel = str_detect(jeopardy$answer, "[aeiou]{3}")) %>%
  filter(vowel == T)
```

```
## # A tibble: 1 x 6
##   category          value question          answer year vowel
##   <chr>          <dbl> <chr>          <chr> <dbl> <lgl>
## 1 LET'S GET SAUCED   1600 "\"The butter of Provence\", this s~ aioli 2008 TRUE
```

Additional Resources

- `stringr` website
- `stringr` cheat sheet
- Regular Expressions cheat sheet
- R for Data Science: Strings